

TD-Shim Introduction

- A lightweight virtual firmware for confidential container

Jiewen Yao, Principal Engineer, Intel Corporation

November, 2021

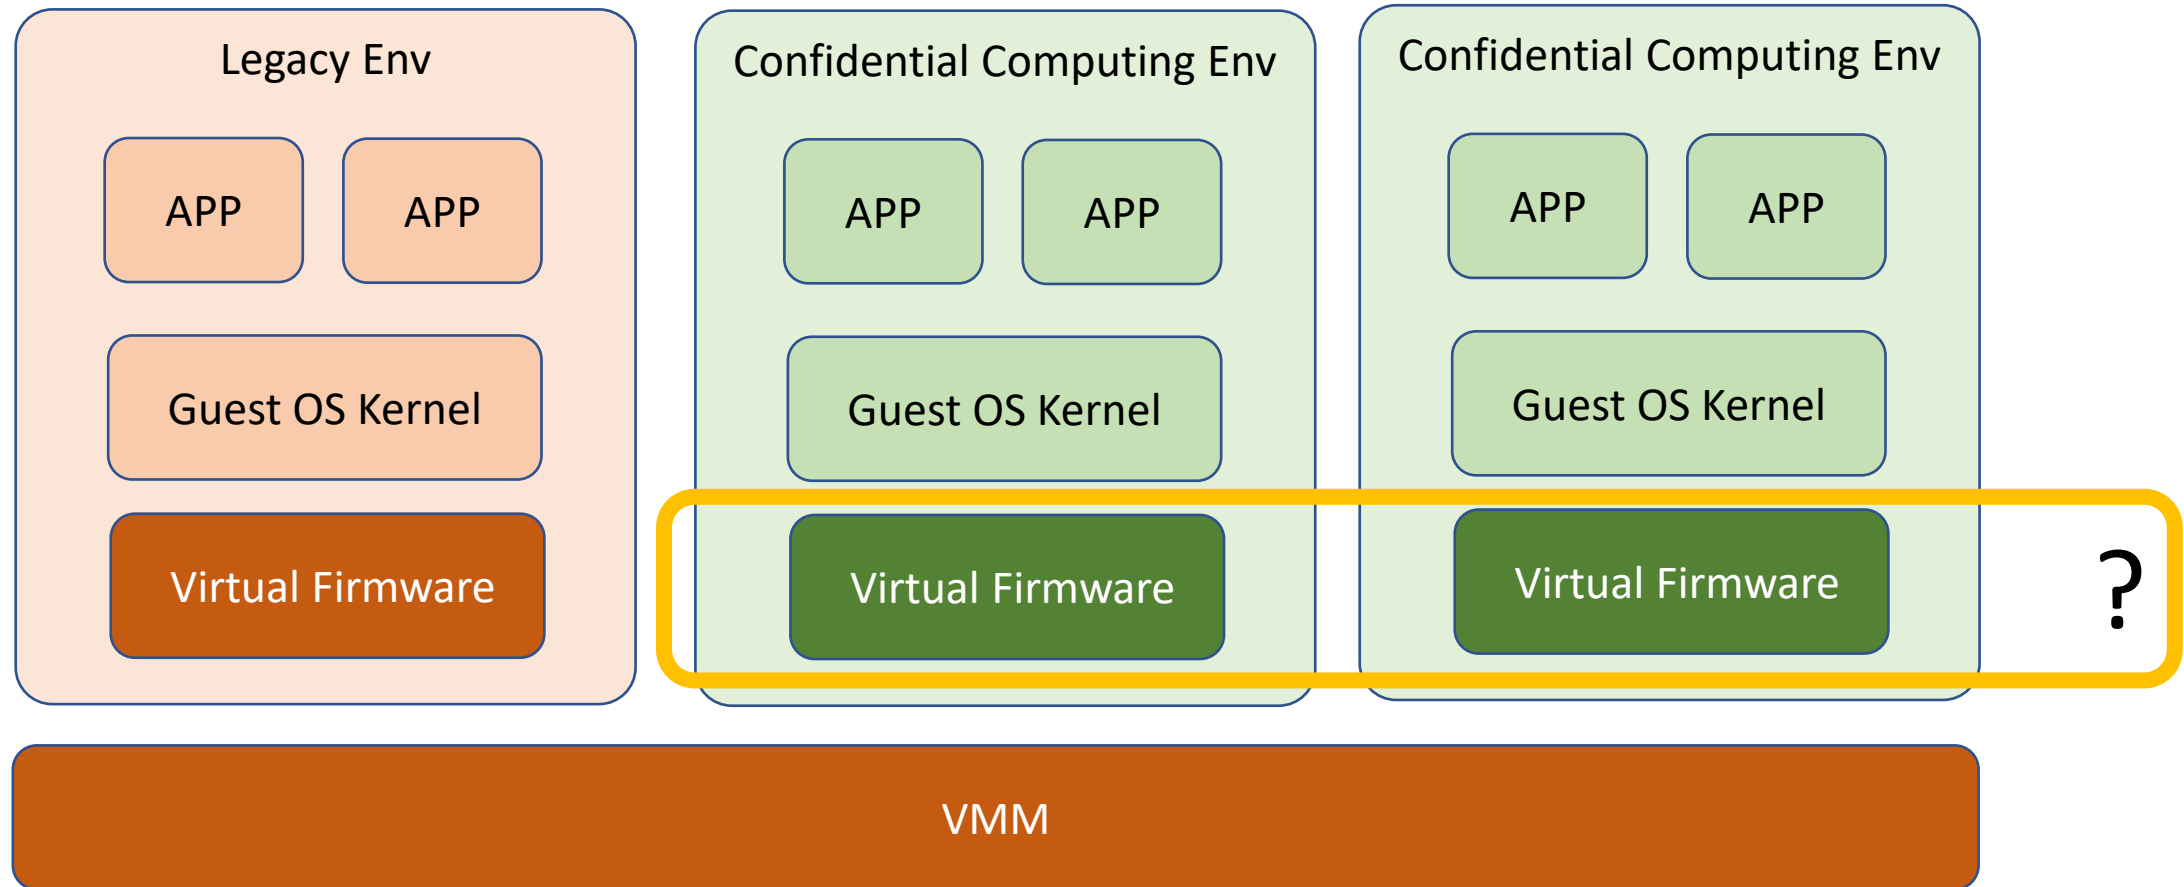
Agenda

- Background
- Why TD-Shim?
- TD-Shim Internal

Background

- Trend for Confidential Computing
 - AMD Secure Encrypted Virtualization (SEV)
 - Intel Trust Domain Extension (TDX)
 - ARM Realm Management Extension (RME)

Need of virtual firmware



Virtual firmware solution today

Main Feature	SeaBIOS	OVMF	cloud-hypervisor-firmware
Hypervisor	XEN, KVM, ...	XEN, KVM, ...	cloud-hypervisor, ...
Arch	16 bit	32bit/64bit	64bit
VMM-BIOS Entrypoint	16bit Reset Vector	16bit Reset Vector	ELF Entrypoint
BIOS-OS Interface	Legacy 16bit INT	UEFI Specification	Linux Boot Protocol
Gap Analysis			
Missing Feature (TDX)	1. Entrypoint – 32bit Reset Vector 2. MP Wakeup – special wakeup structure 3. Memory Initialize – memory accept before use 4. DMA Management – shared/private memory switch 5. Measurement – Runtime Measurement Register (RTMR) extend		
Main Arch Gap (TDX)	16 bit code		ELF Entrypoint No Equivalent Arch

Confidential Computing Solution (TDX)

Main Feature	SeaBIOS ->N/A	OVMF -> TDVF	cloud-hypervisor-firmware -> Td-Shim
Hypervisor	N/A	KVM, cloud-hypervisor,	cloud-hypervisor, KVM, ...
Arch	N/A	64bit	64bit
VMM-BIOS Entrypoint	N/A	Reset Vector	Reset Vector
BIOS-OS Interface	N/A	UEFI Specification, ACPI Specification	Linux Boot Protocol, Executable Payload Boot
Feature	N/A	Support No_CC/TDX/SEV No OVMF arch change	Support TDX only (So far)

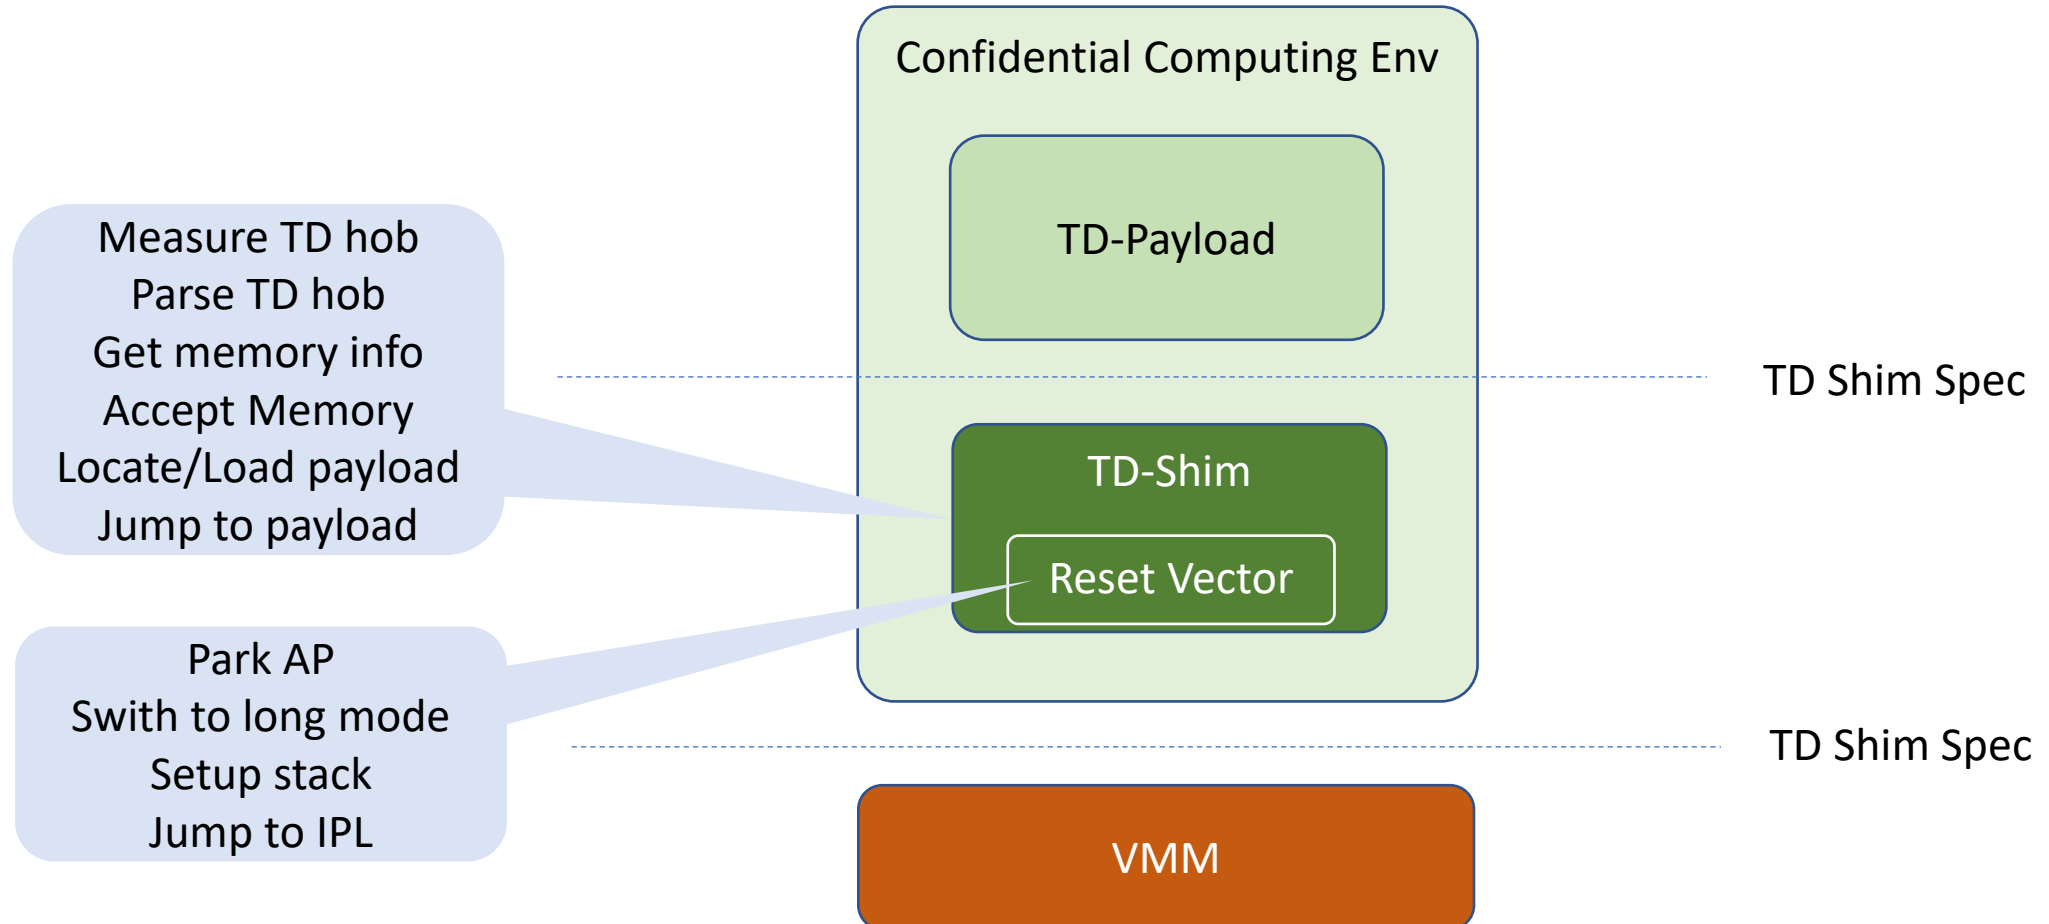
TD-Shim

- A lightweight virtual firmware – bridge the gap between VMM and OS.
- no_std build with Rust.
- Initial version:
 - <https://github.com/confidential-containers/td-shim/pull/2>
 - https://github.com/jyao1/td-shim/tree/init_version
- TD-Shim Spec (0.5):
 - doc/tdshim_spec.md
- TD-Shim Design:
 - doc/design.md

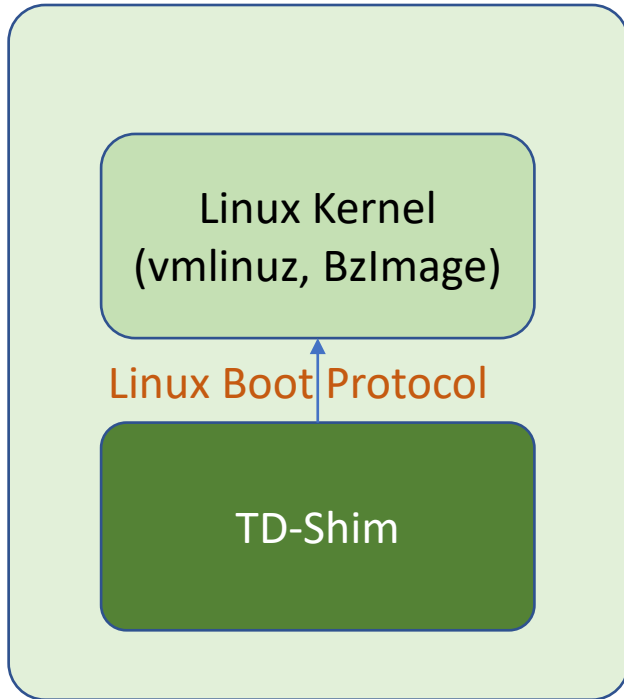
TD-Shim v.s. TDVF(OVMF)

	TD Shim	TDVF
Reset Vector	YES	YES
UEFI Service & Adv Features	NO	YES, Network, File System, etc
OS Runtime	NO	UEFI RT, ACPI ASL
Device Driver	NO	Virtio, PCI, etc
ACPI Table (MP Support)	Static table only (MADT,..). No DSDT.	All (MADT, DSDT, ...)
IRQ Info	Other (Boot Param, ...)	ACPI DSDT
Memory Map	E820 table	UEFI Memory Map
Trusted Boot	YES (RTMR + EventLog)	YES (RTMR + EventLog)
Secure Boot	Optional	Optional (UEFI Secure Boot)
Defense In Depth (DEP, CET)	Default ON	Default OFF
Language	RUST + ASM	C + ASM
Image Size (release)	100K (w/o crypt), 1M (w/ crypt)	4M by default.

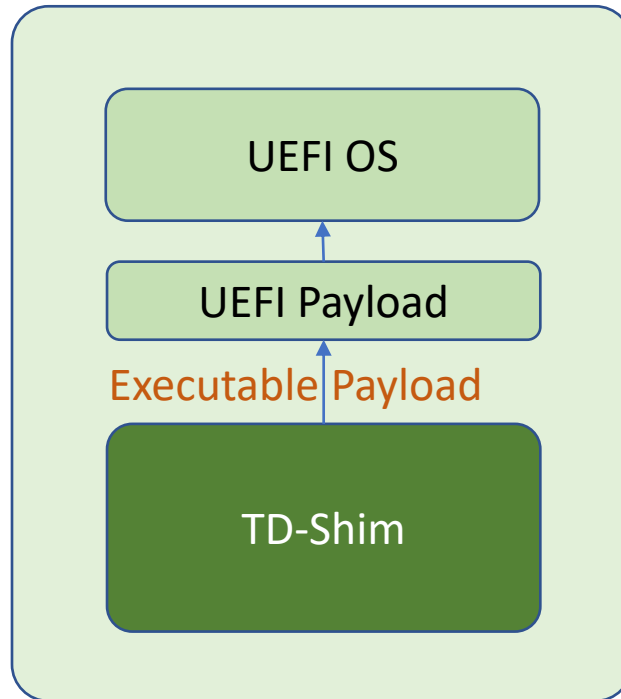
TD-Shim



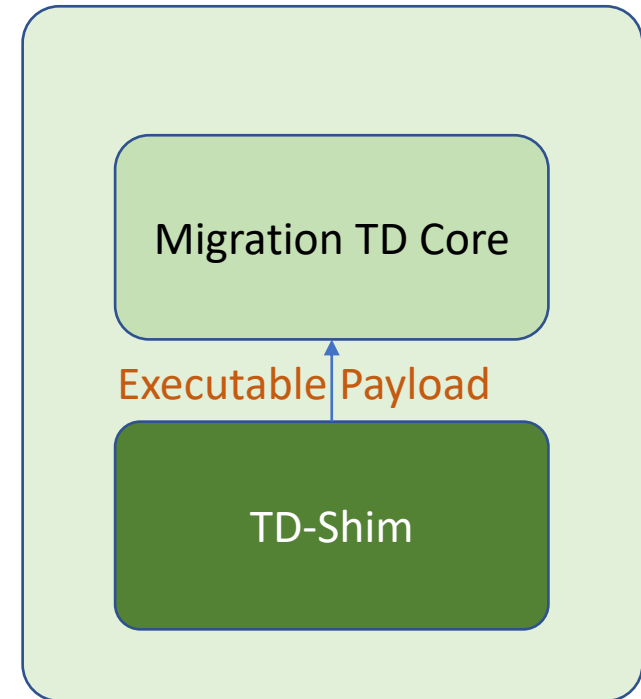
TD-Shim Use Case



Kernel Direct Boot

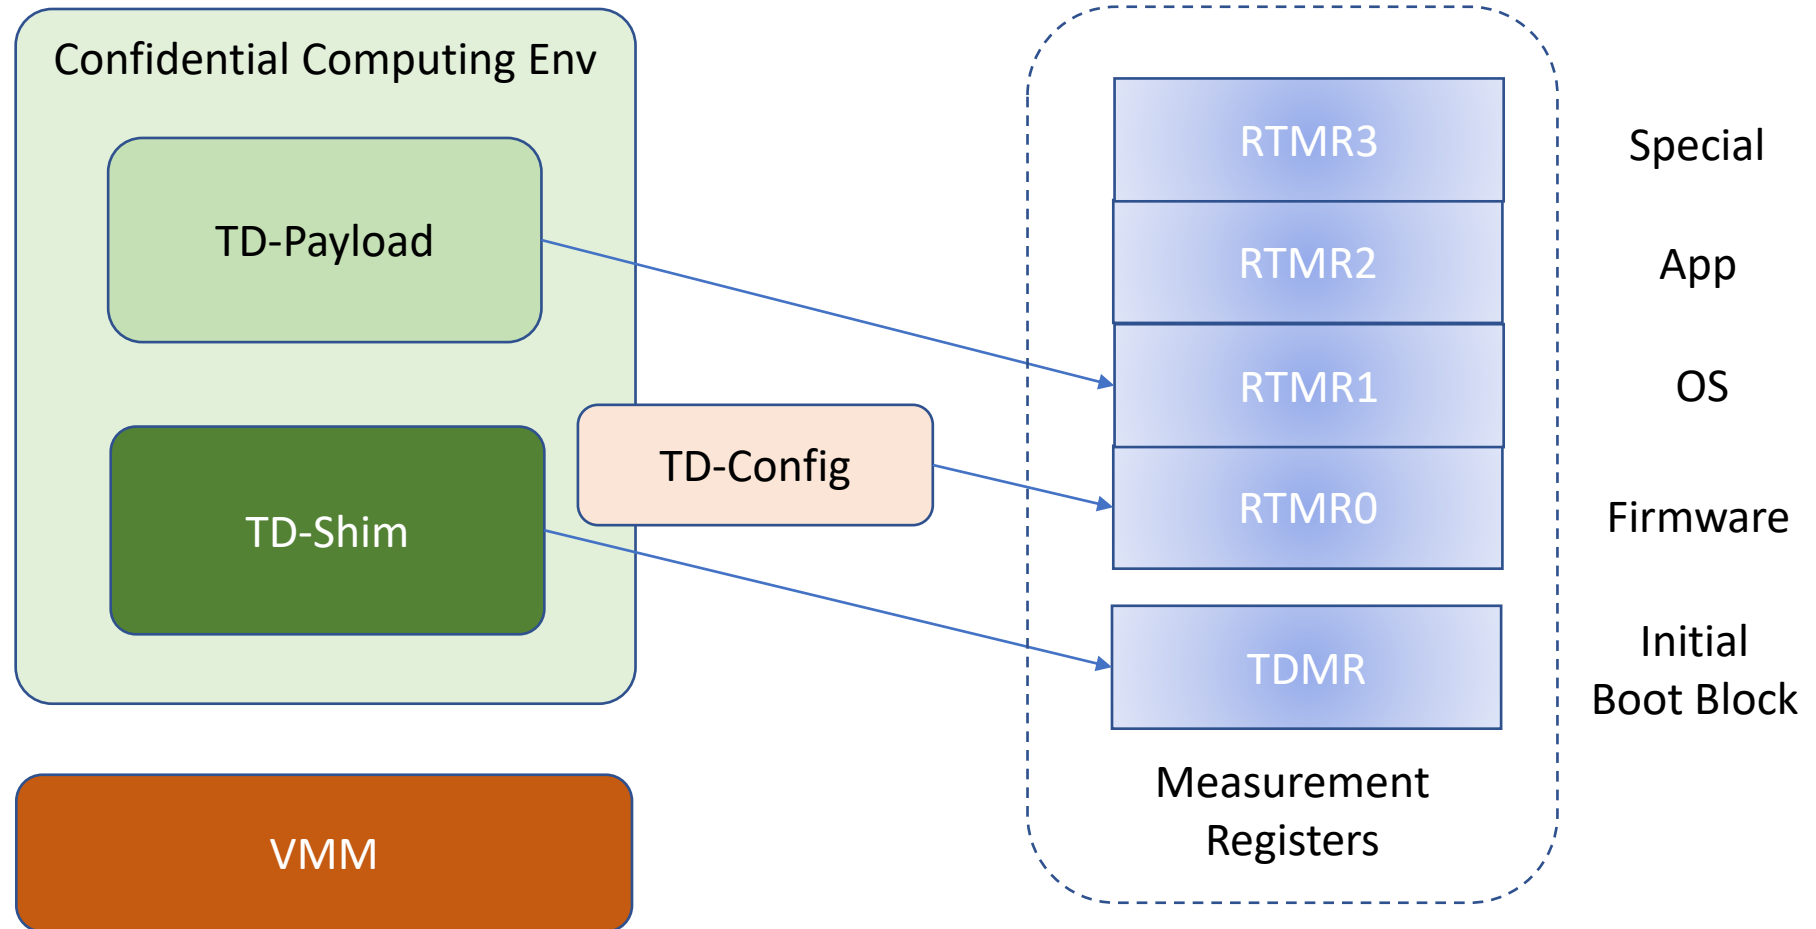


UEFI Boot



Service TD - Migration

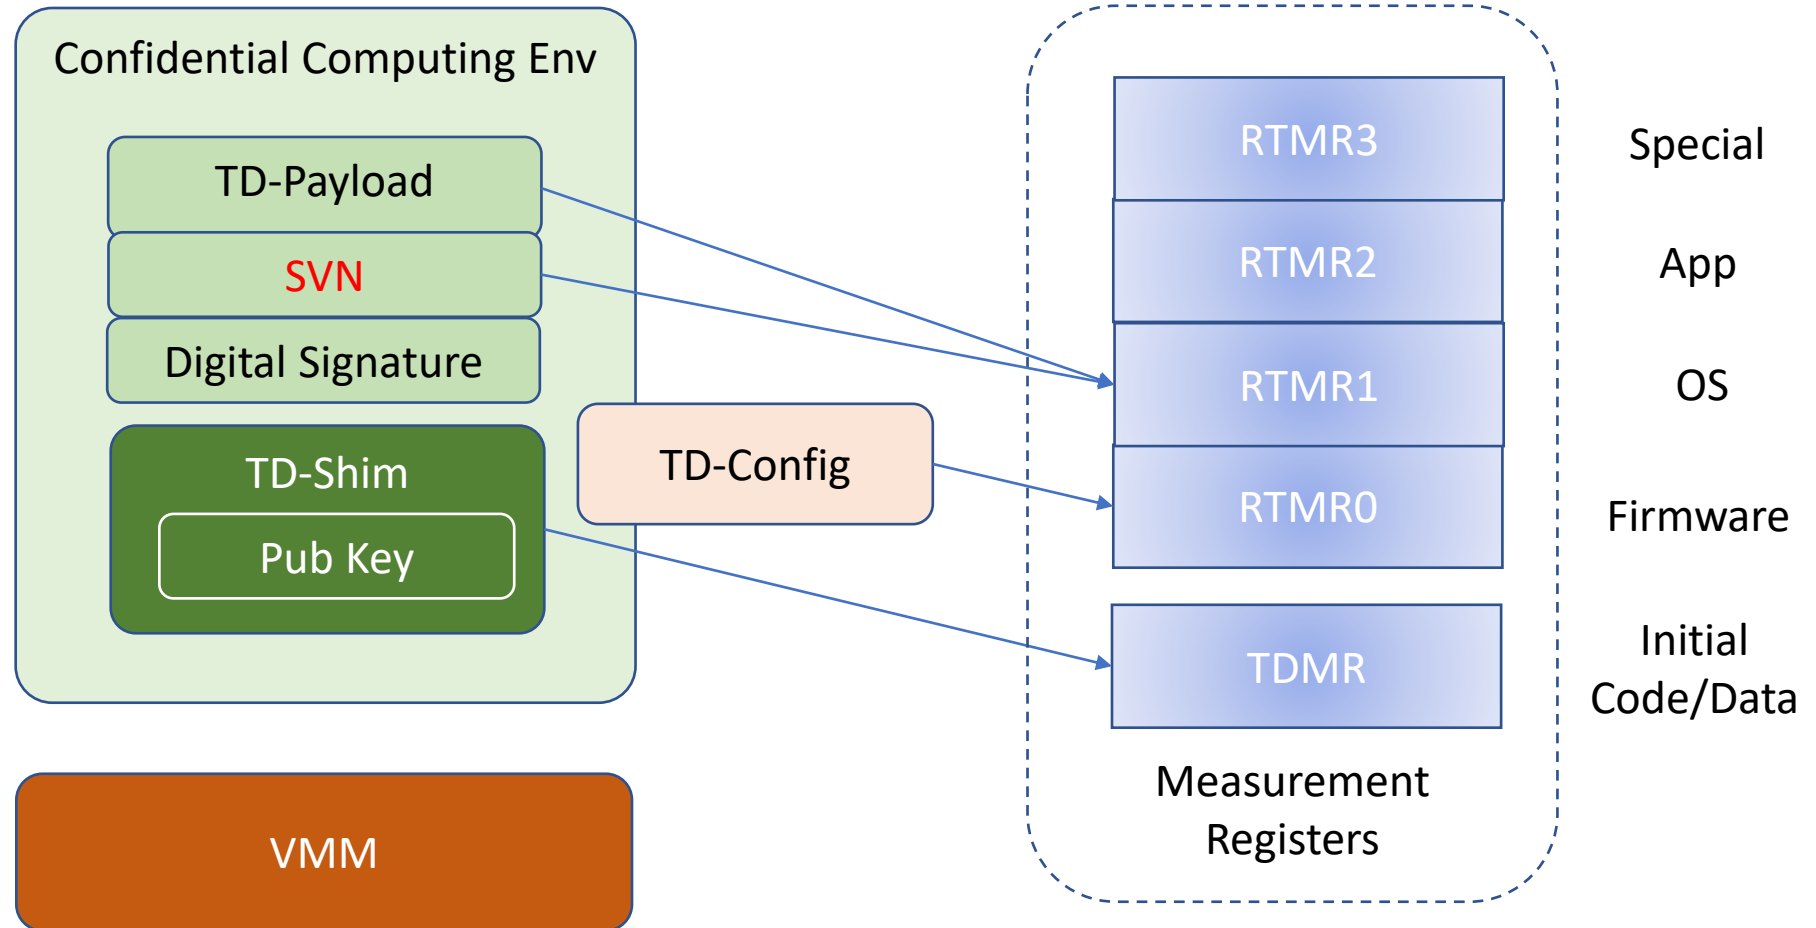
Feature – Trusted Boot & Attestation



Feature – Secure Boot (Optional)

- Use Case - Secure Version Number (SVN) based Attestation
- Problem Statement:
 - How can we do payload attestation based upon SVN, instead of measurement?
 - TD-Shim cannot guarantee the SVN is unmodified.
- Solution:
 - Sign the {payload + SVN}.
 - TD-Shim extends the SVN to RTMR, and creates a dedicate event log.
 - If TD-Shim (in MRTD) is trusted, then SVN (in RTMR) is unmodified.

Feature – Secure Boot with Trusted Boot



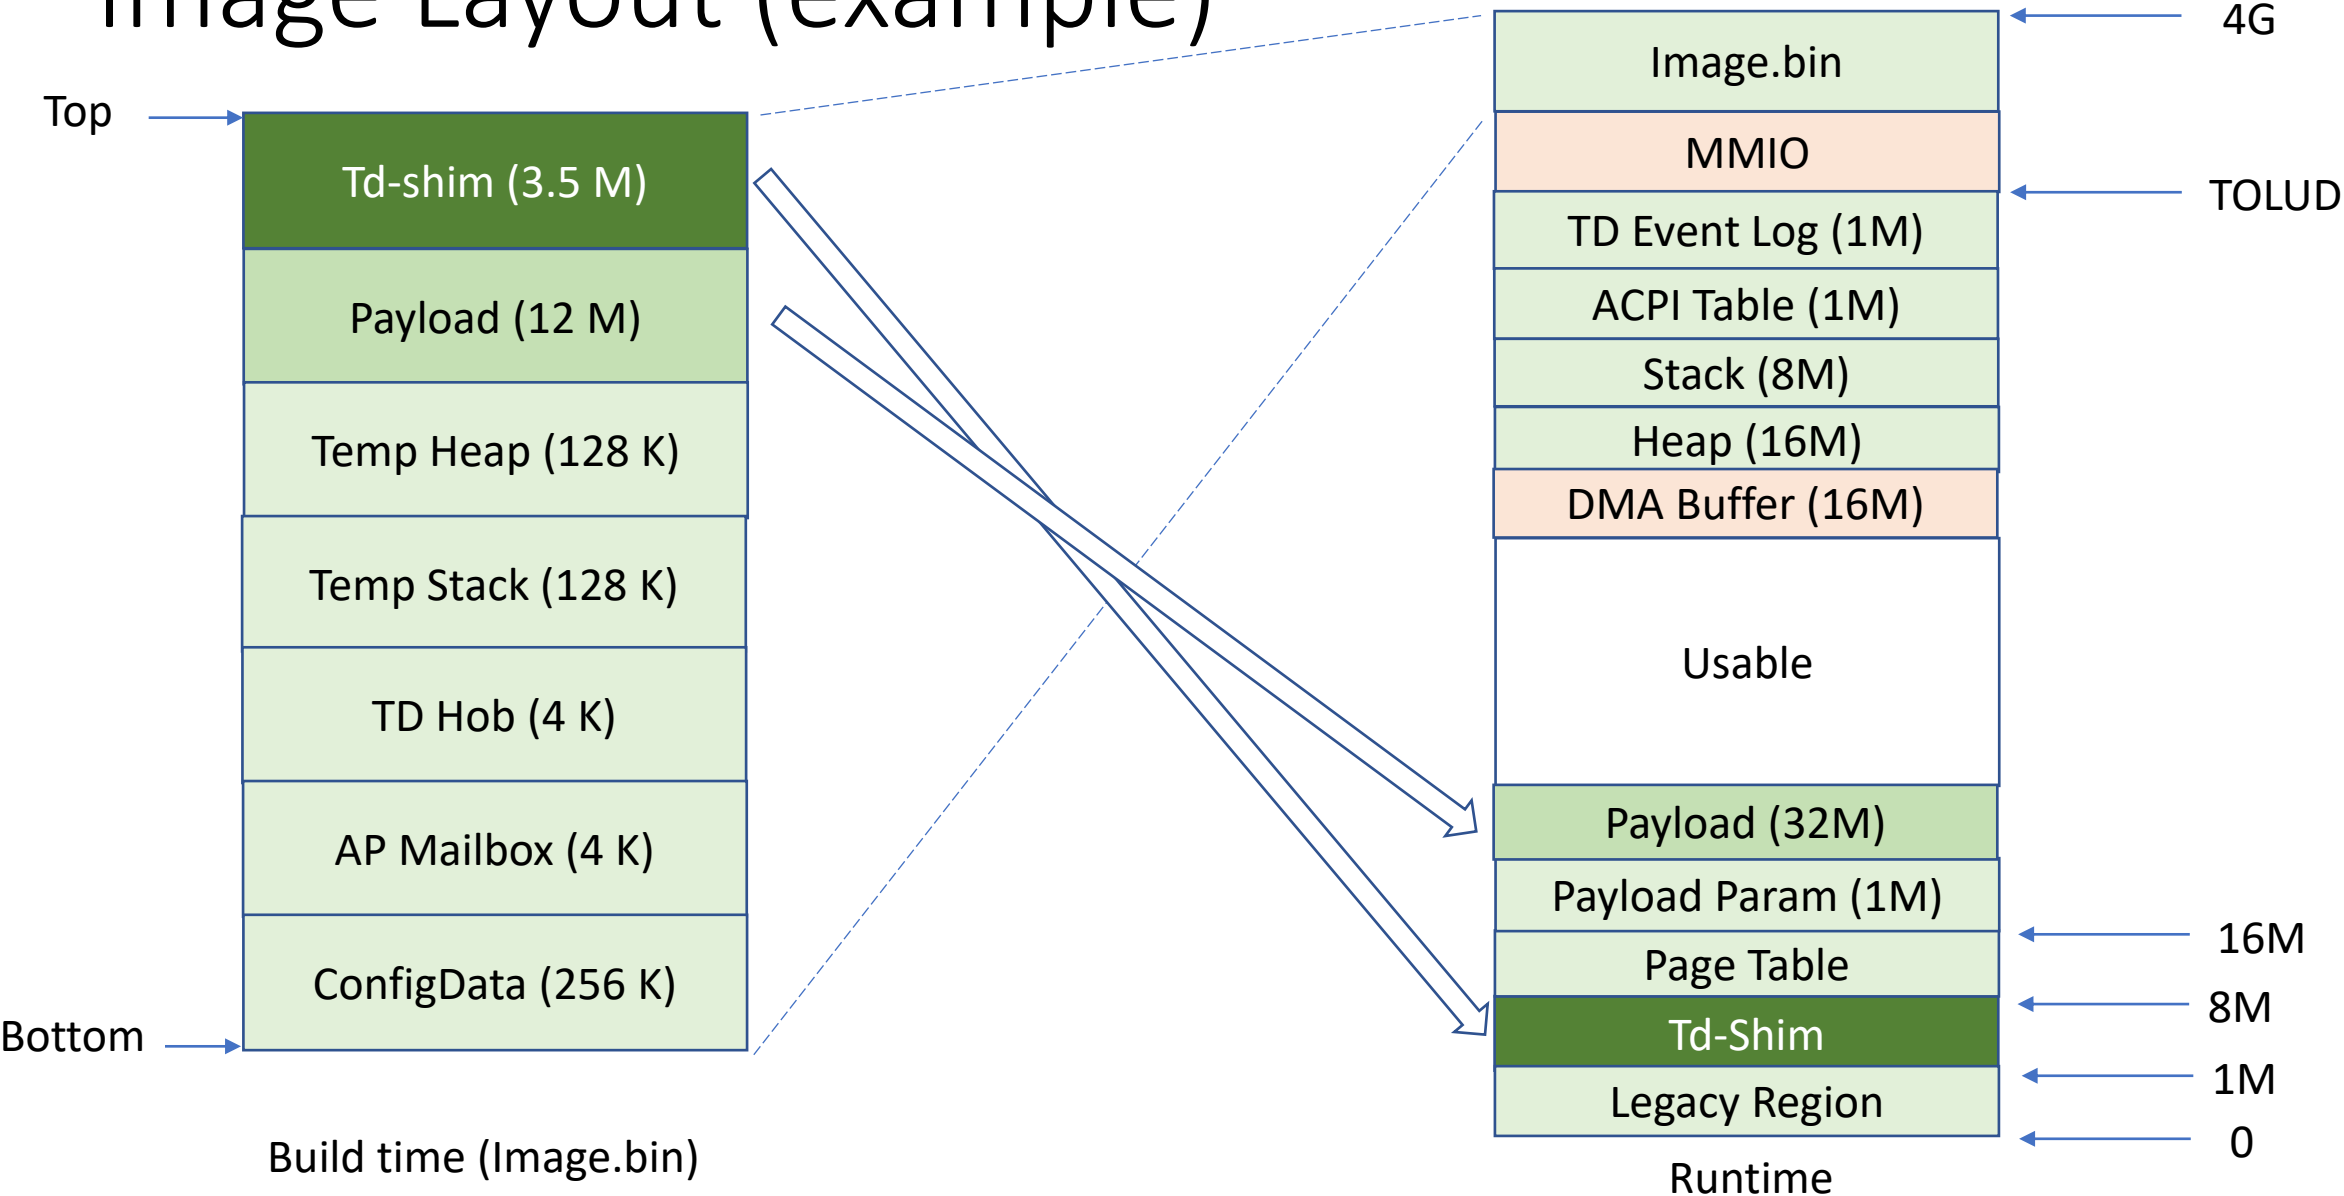
Feature – Defense in depth

- Data Execution Protection (DEP)
 - Page table based enforcement.
 - DataPage = Non-Executable
 - CodePage = Read-Only
- Control Flow Enforcement (CET)
 - Backward-Edge control flow - Shadow Stack (SS)
 - Forward-Edge control flow – Indirect Branch Tracking (IBT)
 - Depend upon compiler (TBD)
- Address Space Layout Randomization (ASLR)
 - Non-Fixed Data Address (Entropy problem, TBD)

Tool

- Everything is self-contained. (No EDKII dependency.)
- **rust-td-tool**: create final binary (reset_vector + td-shim + payload)
- **td-shim**: use x86_64-unknown-uefi (PE)
- **payload**: use x86_64-unknown-uefi (PE) or x86_64-unknown-none (ELF)

Image Layout (example)



Modules

- Td-shim
 - rust-tdshim
 - ResetVector
- TDX related lib
 - tdx-exception
 - tdx-logger
 - tdx-tdcall
- External Dependency
 - crypto - ring (upstream WIP)
- Refer to doc/design.md
- Generic Lib
 - elf-loader
 - pe-loader
 - rust-paging
 - r-uefi-pi (data structure)
 - uefi-pi (fv/hob parsing)
 - fw-pci (for payload)
 - fw-virtio (for payload)
 - fw-vsock (for payload)

Test

- benchmark
 - collect benchmark information, such as stack usage, heap usage, execution time.
- fuzzing-test
 - use afl-fuzz, e.g. pe_loader. (doc/fuzzing.md)
- static code scan
 - run clippy in CI
 - use rudra (doc/rudra.md)
- vulnerable crate scan
 - use cargo-deny (.github/workflows/deny.yml)
- unit test
 - run test in no_std environment. (doc/test_in_no_std.md)
- CI basic test
 - Run test in CI – e.g. clippy (.github/workflows)

Future Plan

- Fix bug & improve quality
 - Improve memory accept performance
 - Feedback is welcome!
- Enable more {hypervisor, guest payload}
 - Feedback is welcome!
- Remove private dependency: crypto – ring
 - Upstream WIP - <https://github.com/briansmith/ring/pull/1406>
- Add std lib support for no_std env (WIP)
 - Evaluating how to add std support to uefi target in rust lang.
 - Important to use other libraries.

Open

- Code Structure
 - Currently, it is flat.
 - Maybe we create some sub folder, to group features?
- Naming
 - <https://github.com/confidential-containers/td-shim/issues/1>
 - Currently, it is **td-shim**.
 - Maybe we rename it to **cc-shim** to allow other arch?